

Multicore Design Considerations

Multicore: The Forefront of Computing Technology

“We’re not going to have faster processors. Instead, making software run faster in the future will mean using parallel programming techniques. This will be a huge shift.”

-- Katherine Yelick, Lawrence Berkeley National Laboratory
from [The Economist: Parallel Bars](#)

- Multicore is a term associated with parallel processing, which refers to the use of simultaneous processors to execute an application or multiple computational threads.
- Parallel programming/processing can be implemented on TI's KeyStone multicore architecture.

Parallel Processing

- Parallel processing divides big applications into smaller applications and distributes tasks across multiple cores.
- The goal is to speed up processing of a computationally-intensive applications.
- Characteristics of computationally-intensive applications:
 - Large amount of data to process
 - Complex algorithms require many computations
- Goals of task partitioning
 - Computational load balancing evenly divides effort among all available cores
 - Minimizes contention of system resources
 - Memory (DDR, shared L2)
 - Transport (Teranet, peripherals)

Parallel Processing: Use Cases



- Network gateway, speech/voice processing
 - Typically hundreds or thousands of channels
 - Each channel consumes about 30 MIPS
- Large, complex, floating point FFT (1M)
- Multiple-size, short FFTs
- Video processing
 - Slice-based encoder
 - Video transcoder (low quality)
 - High-quality decoder

Parallel Processing: Use Cases

- Medical imaging
 - Filtering > reconstruction > post filtering
 - Edge detection
- LTE channel including turbo decoder/encoder
 - Two cores uplink
 - Two cores downlink
- LTE channel excluding turbo decoder
 - Equal to the performance of 30 cores
 - Each core works on a package of bits
- Scientific processing
 - Large complex matrix manipulations
 - Use Case: Oil exploration

Parallel Processing: Control Models

- **Master Slave Model**
 - Multiple speech processing
 - Variable-size, short FFT
 - Video encoder slice processing
 - VLFFT
- **Data Flow Model**
 - High quality video encoder
 - Video decoder
 - Video transcoder
 - LTE physical layer



Parallel Processing: Partitioning Considerations

- **Function driven**
 - Large tasks are divided into function blocks
 - Function blocks are assigned to each core
 - The output of one core is the input of the next core
 - Use cases: H.264 high quality encoding and decoding, LTE
- **Data driven**
 - Large data sets are divided into smaller data sets
 - All cores perform the same process on different blocks of data
 - Use cases: image processing, multi-channel speech processing, sliced-based encoder

Parallel Processing: System Recommendations

- Ability to perform many operations
 - Fixed-point AND floating-point processing
 - SIMD instruction, multicore architecture
- Ability to communicate with the external world
 - Fast two-way peripherals that support high bit-rate traffic
 - Fast response to external events
- Ability to address large external memory
 - Fast and efficient save and retrieve methods
 - Transparent resource sharing between cores
- Efficient communication between cores
 - Synchronization
 - Messaging
 - Data sharing

Parallel Processing: Recommended Tools

- Easy-to-use IDE (Integrated Development Environment)
 - Advanced debug features (system trace, CP tracer)
 - Simultaneous, core-specific debug monitoring
- Real-time operating system (e.g., SYS/BIOS)
- Multicore software development kit
 - Standard APIs simplifies programming
 - Layered abstraction hides physical details from the application
- System optimized capabilities
 - Full-featured compiler, optimizer, linker
 - Third-party support

Example: High Def 1080i60 Video H264 Encoder

- A short introduction to video encoding
 - Pixel format
 - Macroblocks
- Performance numbers and limitations
 - Motion estimation
 - Encoding
 - Entropy encoder
 - Reconstruction
 - Data in and out of the system
 - DDR bandwidth
 - Synchronization, data movement
 - System architecture


Macroblock and Pixel Data

RGB and YUV


$$\begin{bmatrix} Y' \\ U \\ V \end{bmatrix} = \begin{bmatrix} 0.299 & 0.587 & 0.114 \\ -0.14713 & -0.28886 & 0.436 \\ 0.615 & -0.51499 & -0.10001 \end{bmatrix} \begin{bmatrix} R \\ G \\ B \end{bmatrix}$$

$$\begin{bmatrix} R \\ G \\ B \end{bmatrix} = \begin{bmatrix} 1 & 0 & 1.13983 \\ 1 & -0.39465 & -0.58060 \\ 1 & 2.03211 & 0 \end{bmatrix} \begin{bmatrix} Y' \\ U \\ V \end{bmatrix}$$

4:4:4 and 4:2:0 format



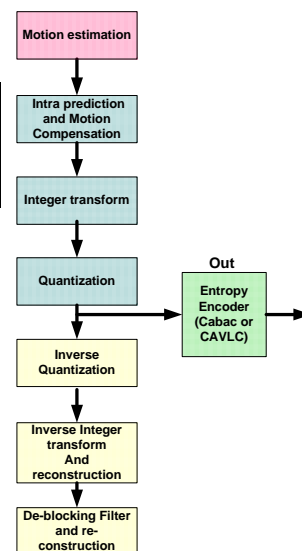
- Typically 8-bit values (10, 12, 14)
- Macroblock = 16x16 pixels



Video Encoder Flow (per Macroblock)

Coder	Width	Height	Frames/Second	MCycles/Second
D1(NTSC)	720	480	30	660
D1 (PAL)	720	576	25	660
720P30	1280	720	30	1850
1080i	1920	1080 (1088)	60 fields	3450

Module	Percentage	Approximate MIPS (1080i)/Second	Number of Cores
Motion Estimation	~50%	1750	2
IP, MC, Transform, Quantization	~12.5%	437.7	0.5
Entropy Encoder	~25%	875	1
IT, IQ and Reconstruction	~12.5%	437.5	0.5



Video Coding Algorithm Limitations

- Motion estimation
 - Depends on the reconstruction of previous (and future) frames
 - Shortcuts can be performed (e.g., first row of frame N does not need last row of frame N-1).
- Intra-prediction
 - Depends on the macroblock above and to the left
 - Must be done consecutively or encoding efficiency is lost (i.e., lower quality for the same number of bits)
- Entropy encoding (CABAC, CAVLC)
 - Must be processed in the macroblock order
 - Each frame is independent of other frames.

How Many Channels Can One C6678 Process?

- Looks like two channels;
Each one uses four cores.
 - Two cores for motion estimation
 - One core for entropy encoding
 - One core for everything else
- What other resources are needed?
 - Streaming data in and out of the system
 - Store and load data to and from DDR
 - Internal bus bandwidth
 - DMA availability
 - Synchronization between cores, especially if trying to minimize delay

What are the System Input Requirements?

- Stream data in and out of the system:
 - Raw data: $1920 * 1080 * 1.5 = 3,110,400$ bytes per frame
= 24.883200 bits per frame (~25M bits per frame)
 - At 30 frames per second, the input is 750 Mbps
 - NOTE: The order of raw data for a frame is Y component first, followed by U and V
- 750 Mbps input requires one of the following:
 - One SRIO lane (5 Gbps raw, about 3.5 Gbps of payload),
 - One PCIe lane (5 Gbps raw)
 - NOTE: KeyStone devices provide four SRIO lanes and two PCIe lanes
- Compressed data (e.g., 10 to 20 Mbps) can use SGMII (10M/100M/1G) or SRIO or PCIe.

How Many Accesses to the DDR?

- For purposes of this example, only consider frame-size accesses.
- All other accesses (ME vectors, parameters, compressed data, etc.) are negligible.
- Requirements for processing a single frame:
 - Retrieving data from peripheral to DDR - 25M bits = 3.125MB
 - Motion estimation phase reads the current frame (only Y) and older Y component of reconstruction frame(s).
 - A good ME algorithm may read up to 6x older frame(s).
 - $7 * 1920 * 1088 = \sim 15\text{M Bytes}$
 - Encoding phase reads the current frame and one old frame. The total size is about 6.25 MB.
 - Reconstruction phase reads one frame and writes one frame. So the total bandwidth is 6.25 MB.
 - Frame compression before or after the entropy encoder is negligible.
 - **Total DDR access for a single frame is less than 32 MB.**

How Does This Access Avoid Contention?

- **Total DDR access for a single frame is less than 32 MB.**
- The total DDR access for 30 frames per second (60 fields) is less than $32 * 30 = 960$ MBps.
- The DDR3 raw bandwidth is more than 10 GBps (1333 MHz clock and 64 bits). 10% utilization reduces contention possibilities.
- DDR3 DMA uses TeraNet with clock/3 and 128 bits. TeraNet bandwidth is $400 \text{ MHz} * 16\text{B} = 6.4$ GBps.

KeyStone SoC Architecture Resources

- 10 EDMA transfer controllers with 144 EDMA channels and 1152 PaRAM (parameter blocks)
 - The EDMA scheme must be designed by the user.
 - The LLD provides easy EDMA usage.
- In addition, Navigator has its own PKTDMA for each master.
- Data in and out of the system (SRIO, PCIe or SGMII) is done using the Navigator.
- All synchronization between cores and moving pointers to data between cores is done using the Navigator.
- IPC provides easy access to the Navigator.

Conclusion

- Two H264 high-quality 1080i encoders can be processed on a single **TMS320C6678**

System Architecture

