

Very Large Fast DFT (VL FFT) Implementation on KeyStone

Multicore Applications

Outlines

- Basic Algorithm for Parallelizing DFT
- Multi-core Implementation of DFT
- Review Benchmark Performance

Goals and Requirements

- **Goal:**
 - To implement very large floating point fast DFT on TI C66x multicore devices
- **Requirements:**
 - FFT sizes: 4K – 1M samples
 - Configurable to run on different number of cores: 1, 2, 4, 8
 - High performance

Algorithm for Very Large DFT

- A generic discrete Fourier transform (DFT) is shown below,

$$y(k) = \sum_{n=0}^{N-1} x(n) e^{j \frac{2\pi}{N} k * n} \quad k = 0, \dots, N-1$$

- Here N is the total size of DFT ,

Algorithm for Very Large DFT

- For very large N , it can be factored into $N = N1 * N2$ and with decimation-in-time, the DFT can be formulated as,

$$\begin{aligned} n1 &= 0, \dots, N1-1, \\ n2 &= 0, \dots, N2-1 \\ n &= n1 * N2 + n2 \end{aligned}$$

$$y(k) = \sum_{n1=0}^{N1-1} \sum_{n2=0}^{N2-1} x(n1 * N2 + n2) e^{j \frac{2\pi}{N1 * N2} k * (n1 * N2 + n2)}$$

$$y(k) = \sum_{n2=0}^{N2-1} \left(\sum_{n1=0}^{N1-1} x(n1 * N2 + n2) e^{j \frac{2\pi}{N1} k * n1} \right) e^{j \frac{2\pi}{N1 * N2} k * n2}$$

$$= \sum_{n2=0}^{N2-1} (DFT_{N1}(k1, n2) e^{j \frac{2\pi}{N1 * N2} k1 * n2}) e^{j \frac{2\pi}{N2} k2 * n2}$$

$$\begin{aligned} k1 &= 0, \dots, N1-1, \\ k2 &= 0, \dots, N2-1 \\ k &= k2 * N1 + k1 \end{aligned}$$

Algorithm for Very Large DFT

- The above DFT formula can be :

$$y(k1 * N2 + k2)$$

$$= \sum_{n2=0}^{N2-1} (DFT_{N1}(k1, n2) e^{j \frac{2\pi}{N1 * N2} k1 * n2}) e^{j \frac{2\pi}{N2} k2 * n2}$$

$$= \sum_{n2=0}^{N2-1} (DFT_{N1}(k1, n2) e^{j \frac{2\pi}{N1 * N2} k1 * n2}) e^{j \frac{2\pi}{N2} k2 * n2}$$

$$= \sum_{n2=0}^{N2-1} A(k1, n2) e^{j \frac{2\pi}{N2} k2 * n2}$$

$$\begin{aligned} k1 &= 0, \dots, N1-1, \\ k2 &= 0, \dots, N2-1 \end{aligned}$$

Algorithm for Very Large DFT

- **A very large DFT of size $N=N1*N2$ can be computed in the following steps:**
 - 1) Formulate input into $N1 \times N2$ matrix
 - 2) Matrix transpose: $N1 \times N2 \rightarrow N2 \times N1$
 - 3) Compute $N2$ FFTs and multiply twiddle factors.
Each FFT is $N1$ size.
 - 4) Matrix transpose: $N2 \times N1 \rightarrow N1 \times N2$
 - 5) Compute $N1$ FFTs. Each is $N2$ size.
 - 6) Matrix transpose: $N1 \times N2 \rightarrow N2 \times N1$

Implementing VLFFT on Multiple Cores

- Two iterations of computations
- 1st iteration
 - $N2$ FFTs are distributed across all the cores.
 - Each core implements matrix transpose and computes **$N2/\text{numCores}$ FFTs and multiplying twiddle factor.**
- 2nd iteration
 - $N1$ FFTs of $N2$ size are distributed across all the cores
 - Each core computes **$N1/\text{numCores}$ FFTs and implements matrix transpose before and after FFT computation.**

Data Buffers

- **DDR3: Three float complex arrays of size N**
 - Input buffer, output buffer, working buffer
- **L2 SRAM:**
 - Two ping-pong buffers, each buffer is the size of 16 FFT input/output
 - Some working buffer
 - Buffers for twiddle factors
 - Twiddle factors for N1 and N2 FFT
 - N2 global twiddle factors

Global Twiddle Factors

- Global Twiddle Factors:

$$e^{j \frac{2\pi}{N1*N2} k1*n2} \quad n2 \in [0, \dots, N2-1] \quad k1 \in [0, \dots, N1-1]$$

- Total of N1*N2 global twiddle factors are required.
- N1 are actually pre-computed and saved.

$$e^{j \frac{2\pi}{N1*N2} n2} \quad n2 \in [0, \dots, N2-1]$$

- The rest are computed during run time.

DMA Scheme

- Each core has dedicated in/out DMA channels
- Each core configures and triggers its own DMA channels for input/output
- On each core, the processing is divided into blocks of 8 FFT each.
- For each block on every core
 - DMA transfer 8 lines of FFT input
 - DSP computes FFT/transpose
 - DMA transfers 8 lines of FFT output

VLFFT Pseudo Code

VLFFT_start:

- 1) **Core0 sends message to each core to start 1st iteration processing.**
- 2) **Each core does the following,**
 - Wait message from core 0 to start,
 - numBlk = 0;
 - While(numBlk < totalBlk)
 - {
 - 1) Trigger DMA to transfer (n+1)th blk from Input Buffer to L2 and to transfer (n-1)th blk output from L2 to Temp Buffer
 - 2) Implement transpose, compute FFT, and multiply twiddle factors for nth blk
 - 3) wait for DMA completion
 - 4) numBlk++
 - }
 - Send a message to core 0
- 3) **Core0 waits for message from each core for completion of its own processing**
- 4) **After receiving all the messages from all the other cores, core0 sends message to each core to start 2nd iteration processing**
- 5) **Each core does the following,**
 - Wait message from core 0 to start,
 - numBlk = 0;
 - While(numBlk < totalBlk)
 - {
 - 1) Trigger DMA to transfer (n+1)th blk from Temp Buffer to L2 and to transfer (n-1)th blk output from L2 to Output Buffer
 - 2) Compute FFT and transpose for nth blk
 - 3) wait for DMA completion
 - 4) numBlk++
 - }
 - Send a message to core 0
- 6) **Core0 waits for message back from each core for completion its own processing**

VLFFT_end:

Matrix Transpose

- The transpose is required for the following matrixes from each core:
 - $N1 \times 8 \rightarrow 8 \times N1$
 - $N2 \times 8 \rightarrow 8 \times N2$
 - $8 \times N2 \rightarrow N2 \times 8$
- DSP computes matrix transpose from L2 SRAM
 - DMA bring samples from DDR to L2 SRAM
 - DSP implements transpose for matrixes in L2 SRAM
 - 32K L1 Cache

Major Kernels

- FFT: single precision floating point FFT from c66x DSPLIB
- Global twiddle factor compute and multiplication: 1 cycle per complex sample
- Transpose: 1 cycle per complex sample

Major Software Tools

- SYS BIOS 6
- CSL for EDMA configuration
- IPC for inter-processor communication

Conclusion

- After the demo ...