Multicore Design Considerations

Multicore: The Forefront of Computing Technology

"We're not going to have faster processors. Instead, making software run faster in the future will mean using parallel programming techniques. This will be a huge shift."
-- Katherine Yelick, Lawrence Berkeley National Laboratory
from The Economist: Parallel Bars

- Multicore is a term associated with parallel processing, which refers to the use of simultaneous processors to execute an application or multiple computational threads.

- Parallel programming/processing can be implemented on TI's KeyStone multicore architecture.

Parallel Processing

- Parallel processing divides big applications into smaller applications and distributes tasks across multiple cores.

- The goal is to speed up processing of a computationally-intensive applications.

- Characteristics of computationally-intensive applications:
  - Large amount of data to process
  - Complex algorithms require many computations

- Goals of task partitioning
  - Computational load balancing evenly divides effort among all available cores
  - Minimizes contention of system resources
    - Memory (DDR, shared L2)
    - Transport (Teranet, peripherals)

Parallel Processing: Use Cases

- Network gateway, speech/voice processing
  - Typically hundreds or thousands of channels
  - Each channel consumes about 30 MIPS

- Large, complex, floating point FFT (1M)

- Multiple-size, short FFTs

- Video processing
  - Slice-based encoder
 - Video transcoder (low quality)
  - High-quality decoder

Parallel Processing: Use Cases

- Medical imaging
  - Filtering > reconstruction > post filtering
  - Edge detection

- LTE channel excluding turbo decoder/encoder
  - Two cores uplink
  - Two cores downlink

- LTE channel including turbo decoder
  - Equal to the performance of 30 cores
  - Each core works on a package of bits

- Scientific processing
  - Large complex matrix manipulations
  - Use Case: Oil exploration

Parallel Processing: Control Models

- ## Master Slave Model
  - – Multiple speech processing
  - – Variable-size, short FFT
  - – Video encoder slice processing
  - – VLFFT

- ## Data Flow Model
  - – High quality video encoder
  - – Video decoder
  - – Video transcoder
  - – LTE physical layer

Parallel Processing: Partitioning Considerations

- **Function driven**
  - Large tasks are divided into function blocks
  - Function blocks are assigned to each core
  - The output of one core is the input of the next core
  - Use cases: H.264 high quality encoding and decoding, LTE
- **Data driven**
  - Large data sets are divided into smaller data sets
  - All cores perform the same process on different blocks of data
  - Use cases: image processing, multi-channel speech processing, sliced-based encoder

Parallel Processing: System Recommendations

- Ability to perform many operations
  - Fixed-point AND floating-point processing
  - SIMD instruction, multicore architecture
- Ability to communicate with the external world
  - Fast two-way peripherals that support high bit-rate traffic
  - Fast response to external events
- Ability to address large external memory
  - Fast and efficient save and retrieve methods
  - Transparent resource sharing between cores
- Efficient communication between cores
  - Synchronization
  - Messaging
  - Data sharing

Parallel Processing: Recommended Tools

- Easy-to-use IDE (Integrated Development Environment)
  - Advanced debug features (system trace, CP tracer)
  - Simultaneous, core-specific debug monitoring
- Real-time operating system (e.g., SYS/BIOS)
- Multicore software development kit
  - Standard APIs simplifies programming
  - Layered abstraction hides physical details from the application
- System optimized capabilities
  - Full-featured compiler, optimizer, linker
  - Third-party support

Example: High Def 1080i60 Video H264 Encoder

- A short introduction to video encoding
    - Pixel format
    - Macroblocks
- Performance numbers and limitations
    - Motion estimation
    - Encoding
    - Entropy encoder
    - Reconstruction
    - Data in and out of the system
    - DDR bandwidth
    - Synchronization, data movement
   - System architecture

Macroblock and Pixel Data

- RGB and YUV

$$\begin{bmatrix} Y' \\ U \\ V \end{bmatrix} = \begin{bmatrix} 0.299 & 0.587 & 0.114 \\ -0.14713 & -0.28886 & 0.436 \\ 0.615 & -0.51499 & -0.10001 \end{bmatrix} \begin{bmatrix} R \\ G \\ B \end{bmatrix}$$

$$\begin{bmatrix} R \\ G \\ B \end{bmatrix} = \begin{bmatrix} 1 & 0 & 1.13983 \\ 1 & -0.39465 & -0.58060 \\ 1 & 2.03211 & 0 \end{bmatrix} \begin{bmatrix} Y' \\ U \\ V \end{bmatrix}$$
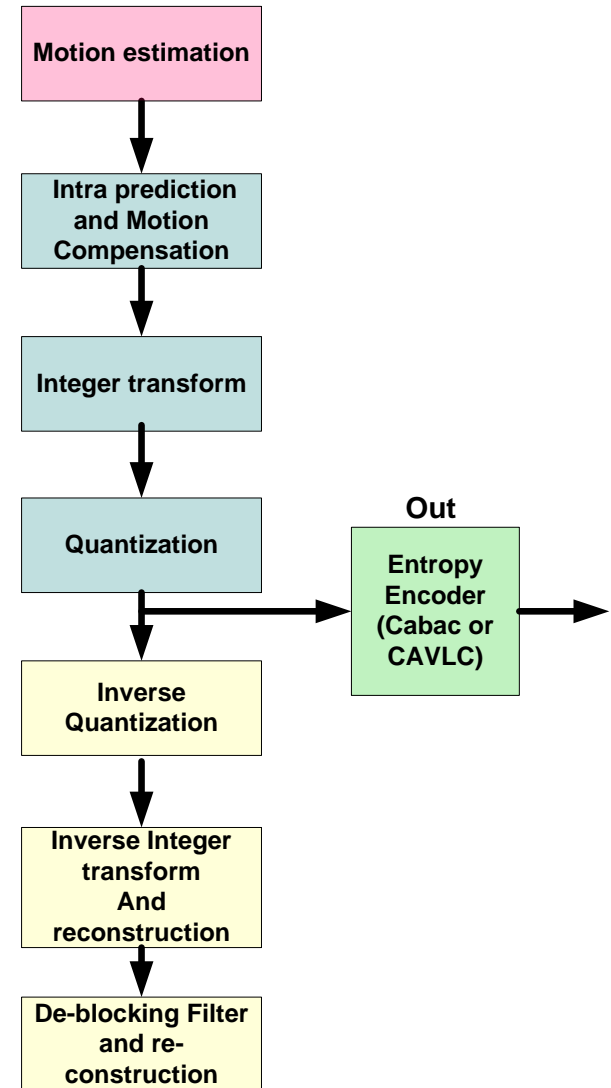
- 4:4:4 and 4:2:0 format

4:4:4   4:2:0

○ -- Pixel with only Y value

◉ -- Pixel with only Cr and Cb values

◉ -- Pixel with Y, Cr and Cb values

- Typically 8-bit values (10, 12, 14)
- Macroblock = 16x16 pixels

Video Encoder Flow (per Macroblock)

Coder	Width	Height	Frames/Second	MCycles/Second
D1(NTSC)	720	480	30	660
D1 (PAL)	720	576	25	660
720P30	1280	720	30	1850
1080i	1920	1080 (1088)	60 fields	3450

Module	Percentage	Approximate MIPS (1080i)/Second	Number of Cores
Motion Estimation	~50%	1750	2
IP, MC, Transform, Quantization	~12.5%	437.7	0.5
Entropy Encoder	~25%	875	1
IT, IQ and Reconstruction	~12.5%	437.5	0.5

Motion estimation

Intra prediction and Motion Compensation

Integer transform

Quantization

Out

Entropy Encoder (Cabac or CAVLC)

Inverse Quantization

Inverse Integer transform And reconstruction

De-blocking Filter and re-construction

Video Coding Algorithm Limitations

- Motion estimation
  - Depends on the reconstruction of previous (and future) frames
  - Shortcuts can be performed (e.g., first row of frame N does not need last row of frame N-1).
- Intra-prediction
  - Depends on the macroblock above and to the left
  - Must be done consecutively or encoding efficiency is lost (i.e., lower quality for the same number of bits)
- Entropy encoding (CABAC, CAVLC)
  - Must be processed in the macroblock order
  - Each frame is independent of other frames.

How Many Channels Can One TMS320C6678 Process?

- Looks like 2 channels; Each one uses 4 cores
  - Two cores for motion estimation
  - One core for entropy encoding
  - One core for everything else
- What other resources are needed?
  - Streaming data in and out of the system
  - Store and load data to and from DDR
  - Internal bus bandwidth
  - DMA availability
  - Synchronization between cores, especially if trying to minimize delay

What are the System Input Requirements?

- Stream data in and out of the system
  - Raw data: 1920 * 1080 * 1.5 = 3,110,400 bytes per frame = 24.883200 bits per frame (~25M bits per frame)
  - At 30 frames per second, the input is 750 Mbps
  - NOTE: The order of raw data for a frame is Y component first, followed by U and V

- 750 Mbps input requires one of the following:
  - One SRIO lane (5 Gbps raw, about 3.5 Gbps of payload),
  - One PCIe lane (5 Gbps raw)
  - NOTE: KeyStone devices provide four SRIO lanes and two PCIe lanes

- Compressed data (e.g., 10 to 20 Mbps) can use SGMII (10M/100M/1G) or SRIO or PCIe.

How Many Accesses to the DDR?

- For purposes of this example, only consider frame-size accesses.
- All other accesses (ME vectors, parameters, compressed data, etc.) are negligible.
- Requirements for processing a single frame:
  - Retrieving data from peripheral to DDR - 25M bits = 3.125MB
  - Motion estimation phase reads the current frame (only Y) and older Y component of reconstruction frame(s).
    - A good ME algorithm may read up to 6x older frame(s).
    - 7 * 1920 * 1088 = ~ 15M Bytes
  - Encoding phase reads the current frame and one old frame. The total size is about 6.25 MB.
  - Reconstruction phase reads one frame and writes one frame. So the total bandwidth is 6.25 MB.
  - Frame compression before or after the entropy encoder is negligible.
  - **Total DDR access for a single frame is less than 32 MB.**

How Does this Access Avoid Contention?

- **Total DDR access for a single frame is less than 32 MB.**

- The total DDR access for 30 frames per second (60 fields) is less than 32 * 30 = 960 MBps.

- The DDR3 raw bandwidth is more than 10 Gbps (1333 MHz clock and 64 bits). 10% utilization reduces contention possibilities.

- DDR3 DMA uses TeraNet with clock/2 and 128 bits. TeraNet bandwidth is 600 MHz * 16B = 9.6 Gbps.

KeyStone SoC Architecture Resources

- 10 EDMA controllers with 144 EDMA channels and 1152 PaRAM (parameter blocks)
  - The EDMA scheme must be designed by the user.
  - The LLD provides easy EDMA usage.
- In addition, Navigator has its own PKTDMA for each master.
- Data in and out of the system (SRIO, PCIe or SGMII) is done using the Navigator.
- All synchronization between cores and moving pointers to data between cores is done using the Navigator.
- IPC provides easy access to the Navigator.

Conclusion

- Two H264 high-quality 1080i encoders can be processed on a single **TMS320C6678**

System Architecture